

Heterogeneity-aware Distributed Parameter Servers

Jiawei Jiang¹,
Bin Cui¹,
Ce Zhang²,
Lele Yu¹,

¹Peking University ²Department of Computer Science

SIGMOD'17, May 14-19, 2017, Chicago, Illinois, USA
Presented by Ching-Yuan, Tsai

Outline

- 1 Introduction
- 2 Preliminaries
- 3 CONSGD
- 4 DYNSGD
- 5 Experiment and Conclusion

Motivation

- The existing systems have bad performance on heterogeneity environments.
 - Network heterogeneity: Data transmission between data centers is normally slower than within data centers.
 - Sharing cluster: Different instances of the same job often have very different execution time.
 - Spot Instances on the Cloud: To minimize cost, a rented cluster often contains different types of instances.

Heterogeneity-aware Distributed Parameter Server

- A novel distributed machine learning system with SGD optimizer.
 - Optimizer: Global learning rate.
 - CONSGD.
 - DYNSGD.
 - Focus on reducing the # updates.

Outline

- 1 Introduction
- 2 Preliminaries**
- 3 CONSGD
- 4 DYNSGD
- 5 Experiment and Conclusion

Distributed SGD

Algorithm 1 SSPSGD [23]

M : # workers, P : # parameter servers, N : # samples, C : # clocks

b : batch size, s : staleness threshold, w_0 : initial parameter

Worker $m = 1, \dots, M$:

- 1: Initialize $w_c^m \leftarrow \text{Pull}(m, 0)$, $c_p \leftarrow 0$
- 2: **for** $c = 0$ to C :
- 3: $u_c^m \leftarrow 0$
- 4: **for** $batch = 1$ to $\frac{N}{bM}$:
- 5: $u_c^m \leftarrow u_c^m - \eta \sum_{i=1}^b \nabla f(x_i, y_i, w_c^m)$
- 6: $w_c^m \leftarrow w_c^m - \eta \sum_{i=1}^b \nabla f(x_i, y_i, w_c^m)$
- 7: $\text{Push}(u_c^m)$
- 8: **if** $c_p < c - s$:
- 9: $(w_{c+1}^m, c_p) \leftarrow \text{Pull}(m, c + 1)$

Parameter Server $p = 1, \dots, P$:

- 1: Initialize $w \leftarrow w_0$, $c_{min} \leftarrow 0$
- 2: **function** $\text{Push}(u_c^m)$:
- 3: $w \leftarrow w + u_c^m$
- 4: **if** all the workers finish c_{min} :
- 5: $c_{min} \leftarrow c_{min} + 1$
- 6: **function** $\text{Pull}(m, c)$:
- 7: **if** $c < c_{min} + s$:

Introduction of existing systems

- BSP system: Bulk Synchronous Parallel.
- ASP system: Asynchronous Parallel.
- SSP system: Stale Synchronous Parallel.

HL

- Decomposing the run time of worker into the computation time t_c^m , the transmission time t_t^m .
- The heterogeneous level of the cluster is measured by the speed gap between the fastest worker and the slowest worker.

$$\bullet \quad HL = \frac{t_c^s + t_t^s}{t_c^f + t_t^f}$$

Experiment

- BSP:Spark, ASP:Petuum, SSP:Bosen.
- They activate the sleep() function in 20%workers to simulate the heterogeneous environment.




Figure 2: Performance of existing systems in the presence of heterogeneity.


Outline

- 1 Introduction
- 2 Preliminaries
- 3 CONSGD**
- 4 DYNSGD
- 5 Experiment and Conclusion


Algorithm

- CONSGD uses a constant global learning rate and multiplies it to each local update.
- $w \leftarrow w + \lambda u_c^m, \lambda \in (0, 1)$

Intuition




(a) A synchronization example



(b) Why SSPSGD cannot work well

Why CONSGD works better?




(c) Why CONSGD works better

Choose λ

- $w_{c+1} = \frac{1}{M} \sum_{i=1}^M (w_c + u_c^i) = w_c + \frac{1}{M} \sum_{i=1}^M u_c^i$
- $\lambda = \frac{1}{M}$ is a good choice.

Limitation



(d) Limitations of CONSGD


Outline

- 1 Introduction
- 2 Preliminaries
- 3 CONSGD
- 4 DYNSGD**
- 5 Experiment and Conclusion


Abstract Model of a Parameter Server

- Denote PS as a tuple $(S, f_{clock}, f_{server}, \lambda, w_0)$, where S is a stream of updates $\{u_1 u_2 \dots\}$.
- $staleness(u_i) = |\{u_j : f_{clock}(j) = f_{clock}(i)\}|$.
- $\lambda(i) = \frac{1}{staleness(i)}$
- $w(PS, i) = w_0 + \sum_j \lambda(j) u_j, \forall f_{clock}(j) < i$.
- $u(PS, i) = \sum_j \lambda(j) u_j, \forall f_{clock}(j) = i$.
 - $w(PS, i + 1) = w(PS, i) + u(PS, i)$

Intuition



(b) Improved calculation graph of DYNSGD



(d) Limitations of CONSGD

Fast staleness computing

- $V(m)$: the version of the local update from worker m .
- $S(v)$: the staleness of the local update stamped with version v .
- $\text{staleness}(u_c^m) = S(V(m))$
- $\Delta u = \frac{1}{d}u_c^m + \frac{d-1}{d}u(PS, v) - u(PS, v) = \frac{1}{d}(u_c^m - u(PC, v))$

Algorithm

Algorithm 2 DYNSGD

c_{min} : clock of the slowest worker, c_{max} : clock of the fastest worker

Parameter Server

- 1: Initialize $c_{min} \leftarrow 0, V \leftarrow 0, S \leftarrow 1$, global parameter $w \leftarrow w_0$
- 2: **function** $Push(u_c^m)$:
- 3: $v \leftarrow V(m)$
- 4: $d \leftarrow S(v)$
- 5: $\Delta u = \frac{1}{d}(u_c^m - u(PS, v))$
- 6: $w \leftarrow w + \Delta u$
- 7: $u(PS, v) \leftarrow u(PS, v) + \Delta u$
- 8: $S(v) \leftarrow S(v) + 1$
- 9: $V(m) \leftarrow V(m) + 1$
- 10: **if** all the items in V is larger than v :
- 11: clear $u(PS, v)$ and $S(v)$
- 12: **if** all the workers finish c_{min} :
- 13: $c_{min} \leftarrow c_{min} + 1$
- 14: **if** $c + 1 > c_{max}$:
- 15: $c_{max} \leftarrow c$
- 16: **function** $Pull(m, c)$:
- 17: **if** $c < c_{min}$:

Partition Synchronization

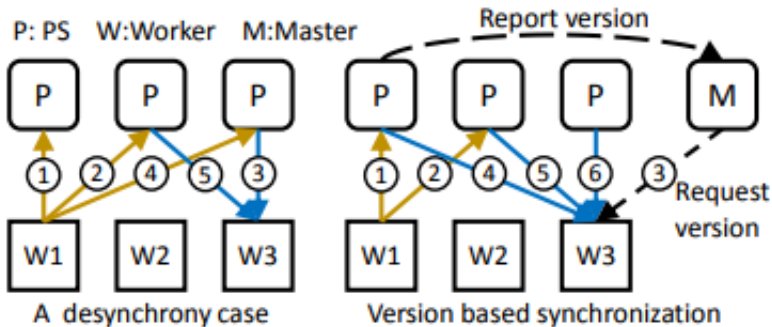


Figure 5: Partition synchronization

Outline

- 1 Introduction
- 2 Preliminaries
- 3 CONSGD
- 4 DYNSGD
- 5 Experiment and Conclusion

Experiment

Metrics	Setting			Spark	Petuum	TF	Petuum	TF	Petuum	CONSGD	DYNSGD	Petuum	CONSGD	DYNSGD
				BSP			ASP		SSP, s=3			SSP, s=10		
run time	LR	URL	HL=1	280	134	172	117	141	87	85	82	78	78	73
			HL=2	555	267	347	625	725	250	196	147	549	116	93
		CTR	HL=1	2352	462	639	748	897	570	456	342	646	532	456
			HL=2	5198	966	1142	4964	5254	2233	1414	546	4501	798	714
	SVM	URL	HL=1	188	84	103	88	122	60	55	51	62	52	44
			HL=2	381	178	226	476	554	188	133	102	454	84	60
		CTR	HL=1	2829	544	660	480	711	528	539	360	544	479	416
			HL=2	6040	1104	1436	6120	6944	3690	1907	873	6534	1036	555
# updates	LR	URL	HL=1	810	810	810	955	937	847	862	833	949	948	887
			HL=2	810	810	810	3858	3901	1243	1062	851	3756	1144	891
		CTR	HL=1	180	180	180	317	303	278	229	164	492	406	353
			HL=2	180	180	180	1964	1936	606	404	204	2035	517	453
	SVM	URL	HL=1	570	570	570	826	848	752	687	628	969	824	692
			HL=2	570	570	570	3318	3276	1128	822	639	3518	986	758
		CTR	HL=1	240	240	240	366	393	322	318	229	467	447	387
			HL=2	240	240	240	2757	2781	890	511	253	2912	773	434
per-update time	LR	URL	HL=1	0.345	0.165	0.212	0.122	0.15	0.103	0.098	0.098	0.082	0.083	0.082
			HL=2	0.691	0.33	0.428	0.162	0.185	0.201	0.184	0.172	0.146	0.102	0.1
		CTR	HL=1	13.0	2.57	3.55	2.36	2.96	2.05	1.99	1.98	1.31	1.3	1.29
			HL=2	28.88	5.37	6.34	2.53	2.71	3.68	3.5	3.33	2.21	1.58	1.54
	SVM	URL	HL=1	0.329	0.148	0.18	0.107	0.144	0.08	0.081	0.08	0.064	0.063	0.061
			HL=2	0.659	0.313	0.396	0.143	0.169	0.166	0.162	0.160	0.129	0.085	0.079
		CTR	HL=1	11.79	2.27	2.75	1.31	1.81	1.64	1.69	1.57	1.16	1.07	1.05
			HL=2	25.17	4.6	5.98	2.22	2.50	4.15	3.72	3.45	2.24	1.34	1.29

Conclusion

- They proposed two heterogeneity-aware distributed SGD algorithms under SSP protocol.

Future Work

- Normalizing updates.